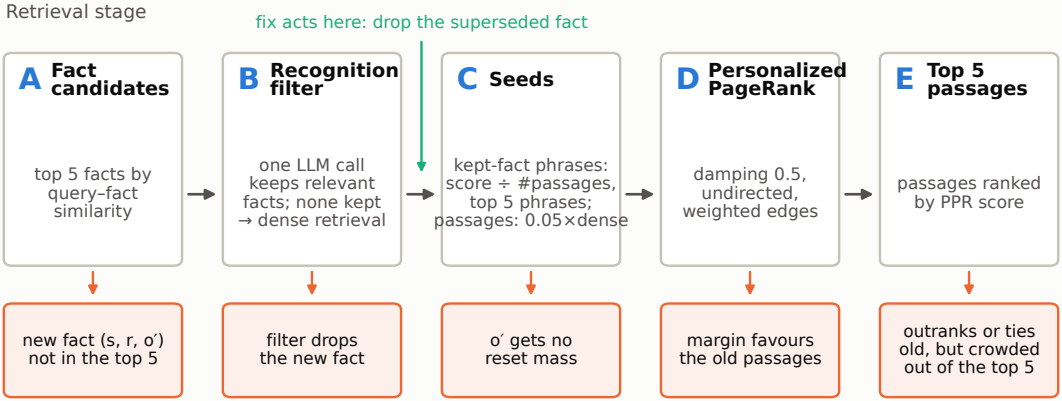
Retrieval stage
fix acts here: drop the superseded fact
A Fact candidates
top 5 facts by query–fact similarity
B Recognition filter
one LLM call keeps relevant facts; none kept → dense retrieval
C Seeds
kept-fact phrases: score ÷ #passages, top 5 phrases; passages: 0.05×dense
D Personalized PageRank
damping 0.5, undirected, weighted edges
E Top 5 passages
passages ranked by PPR score
new fact (s, r, o′) not in the top 5
filter drops the new fact
o′ gets no reset mass
margin favours the old passages
outranks or ties old, but crowded out of the top 5
Failure point for a correction (first-failure attribution, orange)